

University of Pittsburgh
School of Computing and Information

Final Report
INFSCI 2725: Data Analytics
December 4, 2018

Yunshu Liang (yul219)
Qi Lu (qil66)
Erin Price (eep27)
Ziyue Qi (ziq2)
Xiaoqian Xu (xix64)

Table of Contents

Introduction & Problem Definition.....	2
Data Preparation & Preview	3
Feature Analysis	8
Data Cleaning	12
Model Training	21
Evaluate the Model & Cross-Validation.....	24
Results on Kaggle.....	25
Reference.....	26

Introduction & Problem Definition

For the Data Analytics term project, our team named INFSCI 2725 - Fall 2018 includes Yunshu Liang, Qi Lu, Erin Price, Ziyue Qi, and Xiaoqian Xu. The Kaggle project the team will be working on is *Titanic: Machine Learning from Disaster*, located at <https://www.kaggle.com/c/titanic>.

The goal of the project is to predict who did and did not survive on the Titanic, based on data analytics techniques. In this project, we complete the analysis of dataset to predict what sorts of people are likely to survive. In particular, we use Python and implement machine learning algorithms to make a prediction on this binary classification problem.

We solve this problem taking the following steps:

- Define the problem
- Prepare the data & Preview the data
- Feature analysis
- Data cleaning
- Build the model
- Evaluation and Cross-Validation

Data Preparation & Preview

0. Import Libraries

```
import time

import seaborn as sns

import pandas as pd

import matplotlib

import numpy as np

import scipy as sp

import IPython

import sklearn

import matplotlib.pyplot as plt

import itertools

import graphviz

from sklearn.model_selection import RandomizedSearchCV, cross_val_score

from sklearn.ensemble import RandomForestRegressor, RandomForestClassifier

from sklearn import preprocessing, model_selection, ensemble, linear_model, naive_bayes, tree

from math import isnan

from sklearn import datasets

from sklearn.preprocessing import LabelEncoder
```

1. Data Preview

Training and test data are stored in files train.csv and test.csv

1.1. Load train data and test data into pandas data structure—DataFrame and create a dictionary which combined dataTrain and dataTest for the purpose of processing in the same way.

```
Import pandas as pd

csvTrain = "train.csv"
csvTest = "test.csv"
```

```
dataTrain = pd.read_csv(csvTrain)
dataTest = pd.read_csv(csvTest)

dataCleaner = [dataTrain,dataTest]
```

1.2. Preview the data by seeing the first 5 record.

```
print (dataTrain.head())

print (dataTest.head())
```

```
/Users/shine/anaconda/bin/python /Users/shine/Documents/Python/analyze/analyze.py
First 5 records of training data:
PassengerId  Survived  Pclass  \
0            1         0       3
1            2         1       1
2            3         1       3
3            4         1       1
4            5         0       3

                                Name    Sex  Age  SibSp  \
0                                Braund, Mr. Owen Harris    male  22      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38      1
2                                Heikkinen, Miss. Laina  female  26      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)    female  35      1
4                                Allen, Mr. William Henry    male  35      0

Parch  Ticket    Fare  Cabin  Embarked
0      0     A/5 21171   7.2500   NaN      S
1      0      PC 17599  71.2833   C85      C
2      0  STON/O2. 3101282   7.9250   NaN      S
3      0     113803  53.1000  C123      S
4      0    373450   8.0500   NaN      S

First 5 records of test data::
PassengerId  Pclass                                Name    Sex  \
0           892      3                                Kelly, Mr. James    male
1           893      3          Wilkes, Mrs. James (Ellen Needs)  female
2           894      2          Myles, Mr. Thomas Francis    male
3           895      3          Wirz, Mr. Albert    male
4           896      3  Hirvonen, Mrs. Alexander (Helga E Lindqvist)  female

Age  SibSp  Parch  Ticket    Fare  Cabin  Embarked
0  34.5     0     0  330911   7.8292   NaN      Q
1  47.0     1     0  363272   7.0000   NaN      S
2  62.0     0     0  240276   9.6875   NaN      Q
3  27.0     0     0  315154   8.6625   NaN      S
4  22.0     1     1  3101298  12.2875   NaN      S
```

Figure 1 First 5 records

2. Overall Description

2.1. Attributes

```
dataTrain.columns.values
```

```
/Users/shine/anaconda/bin/python /Users/shine/Documents/Python/analyze/analyze.py
Columns:
['PassengerId' 'Survived' 'Pclass' 'Name' 'Sex' 'Age' 'SibSp' 'Parch'
'Ticket' 'Fare' 'Cabin' 'Embarked']
```

Figure 2 Columns

dataTrain.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
PassengerId    891 non-null int64
Survived       891 non-null int64
Pclass         891 non-null int64
Name           891 non-null object
Sex            891 non-null object
Age           714 non-null float64
SibSp          891 non-null int64
Parch          891 non-null int64
Ticket         891 non-null object
Fare           891 non-null float64
Cabin          204 non-null object
Embarked       889 non-null object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.6+ KB
```

Figure 3 Summary of data

From Titanic competition instruction, we know what each column/attribute represents in real world:

PassengerID(int): index of passengers, which might be helpless for our prediction

Survived(int): 0 and 1 representing not survived and survived;

Pclass(int): 1 and 2 and 3. 1 represents highest class;

Name(string): title and name;

Sex(string): female and male;

Age(float): age in years, range from 0.42 to 80;

SibSp(int): of siblings / spouses aboard the Titanic;

Parch(int): of parents / children aboard the Titanic;

Ticket(string): ticket number;

Fare(float): passenger fare;

Cabin(string): cabin number;

Embarked(string): C = Cherbourg, Q = Queenstown, S = Southampton.

2.2. Descriptive statistics of data

`dataTrain.describe()`

Descriptive statistics of training data(numeric):						
	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	
75%	668.500000	1.000000	3.000000	38.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

For strings:						
	Name	Sex	Ticket	Cabin	Embarked	
count	891	891	891	204	889	
unique	891	2	681	147	3	
top	Gill, Mr. John William	male	CA. 2343	G6	S	
freq	1	577	7	4	644	

Figure 4 Descriptive statistics

From statistic summary of the training data, we know number of records, mean, max/min and distribution of each attribute, which influences how we build the model.

2.3. Missing Values

`dataTrain.isnull().sum()`

`dataTest.isnull().sum()`

```
.....
Train columns with null values:
  PassengerId      0
  Survived         0
  Pclass          0
  Name            0
  Sex             0
  Age            177
  SibSp           0
  Parch           0
  Ticket          0
  Fare            0
  Cabin          687
  Embarked        2
dtype: int64
Test/Validation columns with null values:
  PassengerId      0
  Pclass          0
  Name            0
  Sex             0
  Age            86
  SibSp           0
  Parch           0
  Ticket          0
  Fare            1
  Cabin          327
  Embarked        0
dtype: int64
```

Figure 5 Missing Values

Summary:

Age: 117/891 missing in training data, 86/418 missing in test data;

Cabin: 687/891 missing in training data, 327/418 missing in test data;

Fare: 1/418 missing in test data;

Embarked: 2/891 missing in training data.

For further analysis, we'll need a complete dataset without any null values. So missing data should be filled with proper values.

Feature Analysis

1. Sex

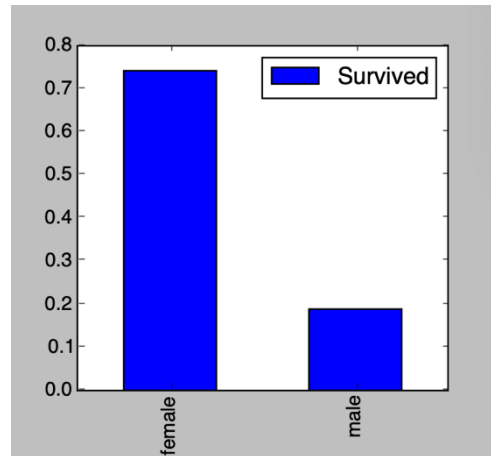


Figure 6 Sex - Survival Rate

Apparently, female has a significantly higher survival rate of roughly 75% than that of male which is below 20%.

2. Pclass

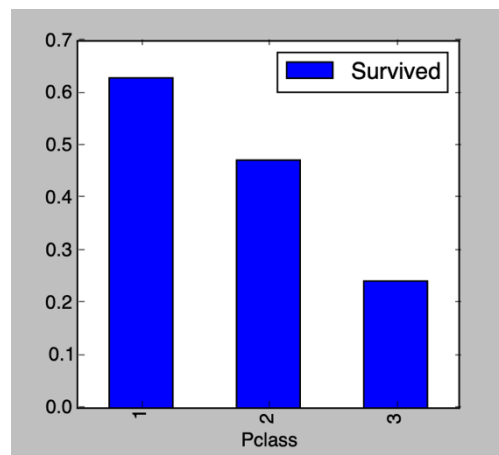


Figure 7 Pclass - Survival Rate

It is clear that first class passengers have the highest survival rate, followed by the second and third class.




Figure 8 Sex & Pclass - Survival Rate

Figure 8 supports the two conclusions we draw above. Female are more likely to survive than male of the same class. And first class has the highest survival rate.

3. Fare

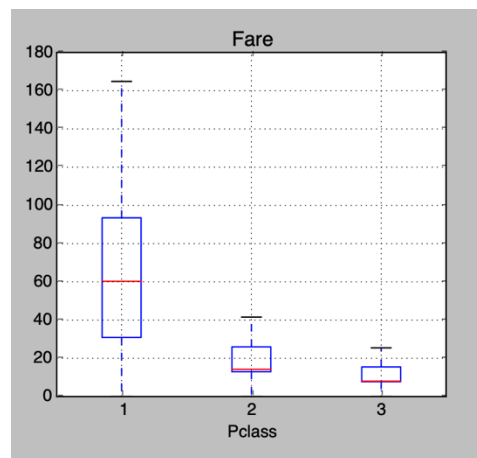


Figure 9 Pclass – Fare

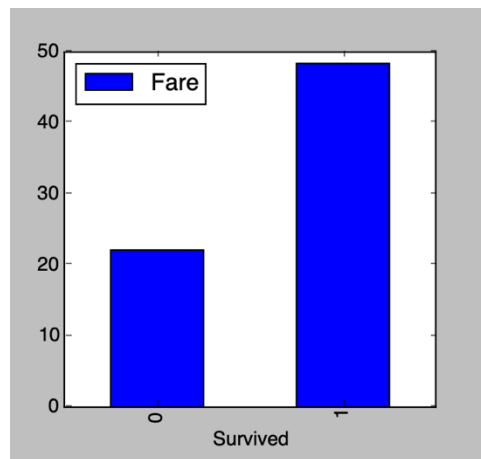


Figure 10 Fare – Survived

We can assume that higher ticket fare represents higher class, leading to the greater chance to survive. Figure 9 and 10 confirms that. The average fare of survived passengers is higher than those who didn't.

4. Age




Figure 31 Age – Survived

From figure 11, it's safe to conclude that young passengers (15-40 years old) get the biggest chance to survive. Also, there's a peak corresponding to infants and young children.

5. Embarked

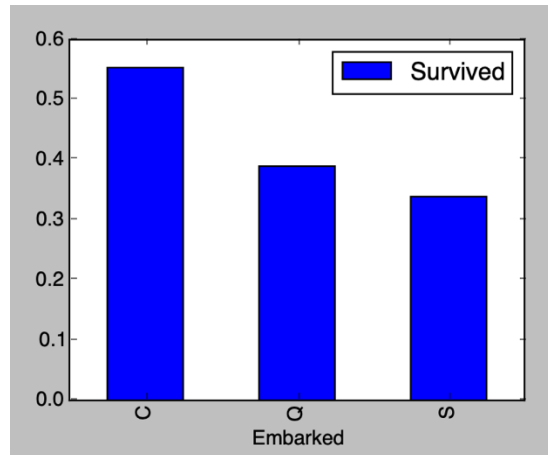


Figure 4 Embarked - Survival Rate

For unclear reasons, departure from port C gives passengers a stronger possibility to survive compared to leaving from the other two ports.

Data Cleaning

1. Completing missing data

1.1 Age

Though we can find some age categories with relatively high survival rate, no clear relation pattern is detected between how exactly age effects the survival rate. Also, the number of missing values are very large. Therefore, we fill the missing Age data in both training and test dataset with predicted values, which are generated by RandomForest algorithm.

```
# Fill missing Age data with predicted values
ageSet = dataTrain[["Age", "Survived", "Fare", "Parch", "SibSp", "Pclass"]]

ageSetNull = ageSet.loc[(dataTrain["Age"].isnull())]

ageSetNotNull = ageSet.loc[(dataTrain["Age"].notnull())]

X = ageSetNotNull.values[:, 1:]

Y = ageSetNotNull.values[:, 0]

RFR = RandomForestRegressor(n_estimators=1000, n_jobs=-1)

RFR.fit(X, Y)

predictAges = RFR.predict(ageSetNull.values[:, 1:])

dataTrain.loc[dataTrain["Age"].isnull(), ["Age"]] = predictAges

dataTrain.info()
```

1.2 Embarked

Only two records have missing Embarked value, so we replace them with mode.

Fill missing Embarked data with mode

```
dataTrain.Embarked[dataTrain.Embarked.isnull()] = dataTrain.Embarked.dropna().mode().values
```

1.3 Cabin

There are 687/891 missing values in training data and 327/418 missing values in test data. The amount of missing values is too large to use either mode or medium to fill with. Therefore, we decide to replace missing Cabin values with "NA".

Fill missing Cabin data with "NA"

```
dataCleaner["Cabin"] = dataCleaner.Cabin.fillna("NA")
```

1.4 Fare

There's only one null value in this column, we fill it in with medium number since this's a numeric attribute.

Fill missing Fare data with "median"

```
dataCleaner['Fare'].fillna(dataCleaner['Fare'].median(), inplace=True)
```

2. Coverting data formats: categorizing continuous variables

2.1 Age

Try three different grouping method and decide on the last one, which is, age 0-15 = group 0, age 15-80 =group 1. This attribute is named AgeBin.

At first, we think we should divide the people based on their escape ability.

```
# first try
# agebin=[0,6,16,35,55,80]
# dataCleaner['AgeBin']=pd.cut(dataCleaner['Age'],bin,labels=[0,1,2,3,4])

# second try
# agebin = [0, 14, 30, 50, 80]
# dataCleaner['AgeBin'] = pd.cut(dataCleaner['Age'], agebin, labels=[0, 1, 2, 3])
```

Then, we found that only kids and teenagers have higher survival rate so that we change another way to divide age bin.




Figure 5 Age - Survival Rate

```
# third try
agebin = [0, 15, 80]
dataCleaner['AgeBin'] = pd.cut(dataCleaner['Age'], agebin, labels=[0, 1])
```

2.2 Fare

Fare is a continuous variable. To turn it into a discrete one, we divide Fare into 3 groups according to the statistical summary showed in figure 3. Values at 25%, 50% and 75% are set as boundaries (Fare ≤ 7.91 , $7.91 < \text{Fare} \leq 14.454$, $14.454 < \text{Fare} \leq 31$, Fare > 31).

0,1,2 represents the three categories mentioned above.

3. Creating columns

3.1 Name

Even though we can't extract much information from the name itself, the title contained in this feature may influence the result. There are 18 titles and we categorize them as 5:

["Capt", "Col", "Major", "Dr", "Rev"] = "Officer",
["Don", "Sir", "the Countess", "Dona", "Lady"] = "Royalty",
["Mme", "Ms", "Mrs"] = "Mrs",
["Mlle", "Miss"] = "Miss",
["Mr"] = "Mr",
["Master", "Jonkheer"] = "Master".




Figure 64 Title - Survival Rate

We can see from figure 14 that females (Mrs./Miss) are still the group that has the highest survival rate. Furthermore, passengers with titles which may represents higher social class (Master/Royalty/Officer) are more likely to survive than average male passengers(Mr.).

Due to the fact that Royalty and Officer are rarely used, we combine these two into one category, "Rare".

1, 2, 3, 4, 5 represents the five categories mentioned above.

3.2 FamilySize

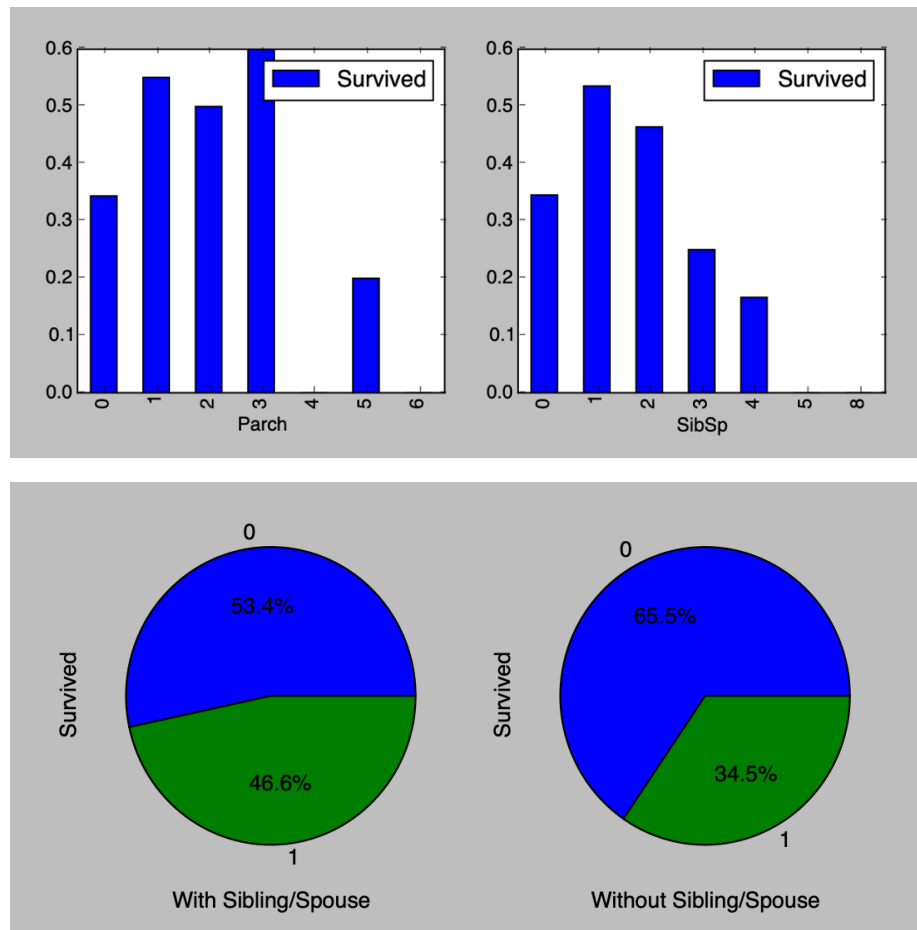


Figure 7 Parch/SibSp - Survival Rate

We can't recognize from figure 15 if Parch or SibSp has impact on survival rate. Considering the fact that these two features both represent the number of family members, we create a new attribute, FamilySize by adding them up.

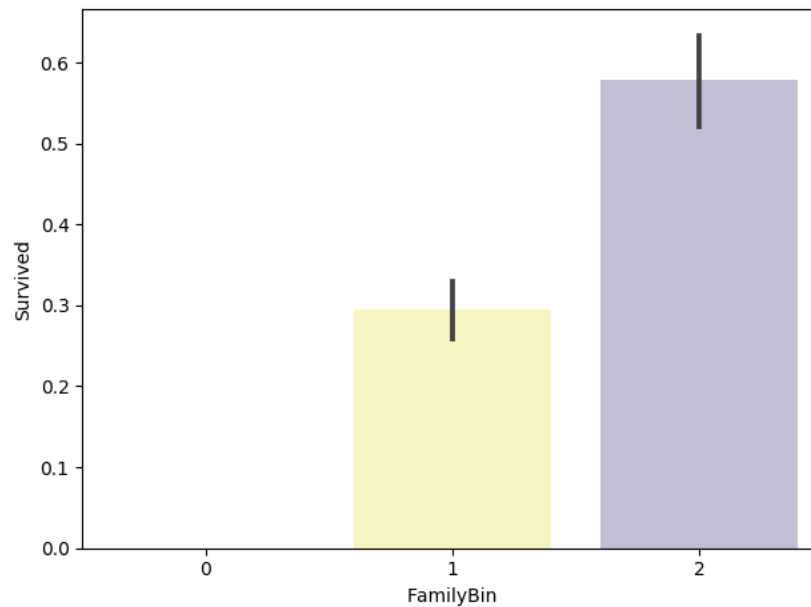


Figure 16 FamilyBin - Survival Rate

From figure 16, we conclude that large families may have difficulty escaping than the small ones. Next, a new feature called FamilyBin is created by dividing FamilySize into 3 categories: 2-4, 4-7 and above 7 persons per family.

2,1,0 represents the three categories mentioned above. Then we can drop columns Parch, SibSp, and FamilySize.

3.3 TicketSize




Figure 17 Ticket - Survival Rate

There are 681 unique Ticket numbers in 891 training data, which means some of the numbers appears more than once. Similar to what we've done to Parch and SibSp, we name the time of a ticket number's occurrences TicketGroup. A new feature called TicketBin is then created by dividing TicketGroup into 3 categories: appear 2-4 times, 4-8 times or once and above 8 times.

2,1,0 represents the three categories mentioned above. Then we can drop columns Ticket and TicketSize.

4. Correcting: Dealing with outliers

A special case is taken into consideration. We divide the passengers into different groups based on their surname. We divide the groups which have more than one passenger into two sets: female and children, male adult. We can see that female and children group mostly all died or all survived. It is the same with male adult group. We define the group in female and children group with 0.0 survival rate as died group. Then, we define the group in male adult group with 1.0 survival rate as survived group. We need to deal with variables in these two abnormal groups by changing their values and attributes. For example, we can change the sex in the survived group from 'female' to 'male'.

5. Dropping: PassengerID & Cabin

Unrelated columns, drop.

6. Clean data

new schema of dataset

print(train.info())

print(test.info())

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 891 entries, 0 to 890
Data columns (total 9 columns):
Embarked      891 non-null int64
Fare           891 non-null int64
Pclass        891 non-null int64
Sex            891 non-null int64
Survived       891 non-null float64
Title         891 non-null int64
FamilyBin     891 non-null int64
TicketBin     891 non-null int64
AgeBin        891 non-null category
dtypes: category(1), float64(1), int64(7)
memory usage: 63.5 KB
None
<class 'pandas.core.frame.DataFrame'>
Int64Index: 418 entries, 891 to 1308
Data columns (total 8 columns):
Embarked      418 non-null int64
Fare           418 non-null int64
Pclass        418 non-null int64
Sex            418 non-null int64
Title         418 non-null int64
FamilyBin     418 non-null int64
TicketBin     418 non-null int64
AgeBin        418 non-null category
dtypes: category(1), int64(7)
memory usage: 26.5 KB
None
```

Figure 18 Clean data schema

Model Training

1. Model Selection

We have tried multiple machine learning algorithms and compared them for different scenarios. We choose:

- 1) Ensemble Methods: RandomForest Classifier, ExtraTrees Classifier and AdaBoost Classifier
- 2) Navies Bayes: naive_bayes.BernoulliNB, naive_bayes.GaussianNB()
- 3) DecisionTree: tree.DecisionTreeClassifier(), tree.ExtraTreeClassifier()
- 4) Logistic Reegression

	MLA Name	MLA Test Accuracy Mean
2	RandomForestClassifier	0.822761
1	ExtraTreesClassifier	0.822015
6	DecisionTreeClassifier	0.81791
7	ExtraTreeClassifier	0.815299
3	LogisticRegressionCV	0.808209
0	AdaBoostClassifier	0.801866
4	BernoulliNB	0.78806
5	GaussianNB	0.758209

Figure 19 Machine learning algorithms -test accuracy




Figure 20 Machine learning algorithms - Test accuracy score

We can learn from figure 20 that Random Forest Model have highest score so that we decide to use it and tune this model.

2. Tune Random Forest

Originally use parameters as follows:

```
{'bootstrap': True,  
 'class_weight': None,  
 'criterion': 'gini',  
 'max_depth': None,  
 'max_features': 'auto',  
 'max_leaf_nodes': None,  
 'min_impurity_decrease': 0.0,  
 'min_impurity_split': None,  
 'min_samples_leaf': 1,  
 'min_samples_split': 2,  
 'min_weight_fraction_leaf': 0.0,  
 'n_estimators': 'warn',  
 'n_jobs': None,  
 'oob_score': False,  
 'random_state': 0,  
 'verbose': 0,  
 'warm_start': False}
```

Then tune the parameters with RandomizedSearchCV and find the best parameters:

```
{'bootstrap': False,  
 'class_weight': None,  
 'criterion': 'gini',  
 'max_depth': 3,  
 'max_features': 'auto',
```

```
'max_leaf_nodes': None,  
'min_impurity_decrease': 0.0,  
'min_impurity_split': None,  
'min_samples_leaf': 2,  
'min_samples_split': 2,  
'min_weight_fraction_leaf': 0.0,  
'n_estimators': 1000,  
'n_jobs': None,  
'oob_score': False,  
'random_state': 0,  
'verbose': 0,  
'warm_start': False}
```


Evaluate the Model & Cross-Validation

1. Cross-Validation

We get 83.27% mean CV score using Random Forest Model.

#Cross Vaildation

```
cv_score = cross_val_score(RFC1,train[dataTrain_x_bin], train[Target], cv= 10)

print("CV Score : Mean - %.7g | Std - %.7g " % (np.mean(cv_score), np.std(cv_score)))

CV Score : Mean - 0.8327554 | Std - 0.04070334
```

Figure 21 Cross Validation

2. Evaluate the model

We also evaluate another model using decision-tree model and achieve the accuracy of 77.9%. So we believe that the first model is better than decision-tree model.

3. Submit

```
submit = dataTest[['PassengerId','Survived']]

submit.to_csv("submit.csv", index=False)

print('Validation Data Distribution: \n', dataTest['Survived'].value_counts(normalize = True))

submit.sample(10)

print("Done")
```

Results on Kaggle

1 Active Competition



Titanic: Machine Learning from Disaster

Start here! Predict survival on the Titanic and get familiar with ML basics
Getting Started · Ongoing · tutorial, tabular data, binary classification



345/10628
Top 4%

We submitted to Kaggle 22 times in order to see our result and improve our model.

Accuracy/Score: 0.83253

Your most recent submission

Name	Submitted	Wait time	Execution time	Score
submit.csv	an hour ago	0 seconds	3 seconds	0.83253
Complete				

Rank: 345/10628 (Top 4%)

345	new	INFSCI 2725 - Fall 2018		0.83253	22	1h
Your Best Entry ↑						
Your submission scored 0.83253, which is not an improvement of your best score. Keep trying!						

Leaderboard

Overview	Data	Kernels	Discussion	Leaderboard	Rules	Team	My Submissions	Submit Predictions
339	▲ 266	zaid34						0.83253 45 2d
340	new	Dhiego Souto						0.83253 18 2d
341	new	lumimevi						0.83253 4 1d
342	▲ 2295	Pioneer						0.83253 13 7h
343	new	FawnyLu						0.83253 11 2h
344	new	Joe_						0.83253 4 3h
345	new	INFSCI 2725 - Fall 2018						0.83253 22 1h
Your Best Entry ↑								
Your submission scored 0.83253, which is not an improvement of your best score. Keep trying!								
346	new	yangyin						0.82775 4 2mo
347	▼ 44	universehiro						0.82775 15 1mo
348	▼ 44	Joao Megale						0.82775 10 7d
349	new	hanhua						0.82775 3 2mo

Reference

```
1 # -----  
2 # Titanic: Machine Learning from Disaster  
3 # Kaggle Competition  
4 # -----  
5 # INFSCI 2725: Data Analytics  
6 # Fall 2018  
7 # -----  
8 # Yunshu Liang (yul219)  
9 # Qi Lu (qil66)  
10 # Erin Price (eep27)  
11 # Ziyue Qi (ziq2)  
12 # Xiaoqian Xu (xix64)  
13 # -----  
14 # Reference: https://zhuanlan.zhihu.com/p/31743196  
15 # https://www.kaggle.com/ldfreeman3/a-data-science-framework-to-achieve-99-accuracy  
16 # https://zhuanlan.zhihu.com/p/33733586  
17 # -----
```

Our code includes our references in the comments at the beginning of the file. We also list the references here:

- <https://www.kaggle.com/ldfreeman3/a-data-science-framework-to-achieve-99-accuracy>
- <https://zhuanlan.zhihu.com/p/31743196>
- <https://zhuanlan.zhihu.com/p/33733586>